

Multithreading Basics and Programming

Prof. Yung-Pin Cheng

2017 03

History

- In the past (< 1996), when there was no multithreading support from O.S., how a server (such as telnet, BBS server) is implemented?
- The era of multi-process programming

Unix's **fork()** revisited

Process management

Call	Description
pid = fork()	Create a child process identical to the parent
pid = waitpid(pid, &statloc, options)	Wait for a child to terminate
s = execve(name, argv, environp)	Replace a process' core image
exit(status)	Terminate process execution and return status

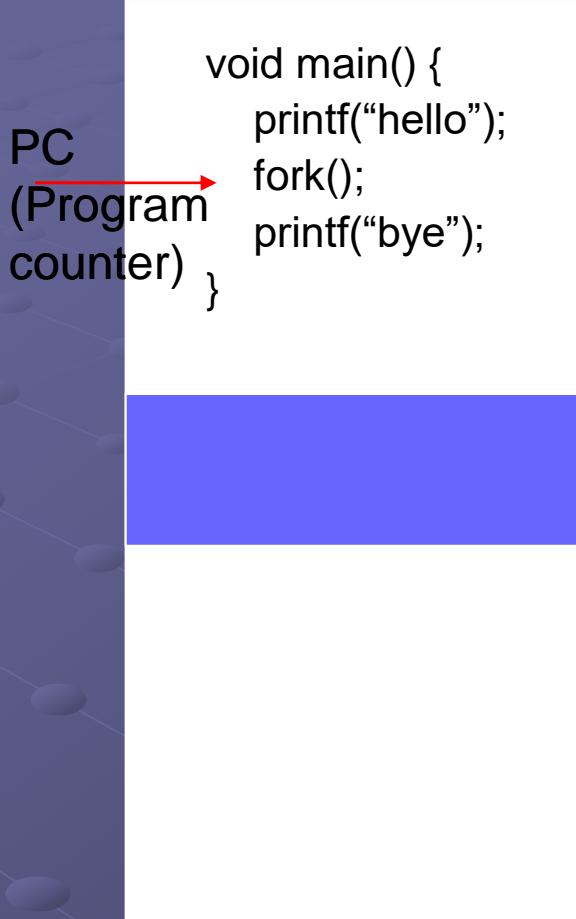
An example of Unix fork()

```
void main() {  
    printf("hello");  
    fork();  
    printf("bye");  
}
```

The output

hello
bye
bye

How `fork()` is implemented in Unix?



Process's Image in Memory

fork() example again

```
void main() {  
    if (fork() == 0)  
        printf(" in the child process");  
    else  
        printf(" in the parent process");  
}
```

An old-fashioned concurrent server

```
main() {
// create a TCP/IP socket to use
s =socket(PF_INET,SOCK_STREAM ,0);

// bind the server address
Z = bind(s, (struct sockaddr *)&adr_srvr,
len_inet);

// make it a listening socket
Z = listen(s, 10);

// start the server loop
for (;;) {
    // wait for a connect
    c = accept(s, (struct sockaddr*)
&adr_clnt,& len_int);
```

```
PID = fork();
if (PID > 0) {
    // parent process
    close(c);
    continue ;
}
// child process
rx = fdopen(c,"r");
tx = fdopen(dup(c), "w");

// process client's request
.....
fclose(tx); fclose(rx);
exit(0);
```

Problems with old-fashioned multi-processes programming

- Context switching overhead is high
- Communications between forked processes also has high cost
 - typical mechanism – via **IPC (interprocess communication)** such as **shared memory, semaphore, message queue**
 - these are cross address space communication
 - invoking of IPC cause mode switch – **may cause a process to block.**

The concept of a *process*

- However, processes are still the very basic element in O.S.
- UNIT of resources ownership
 - Allocated with virtual address space + control of other resources such as I/O, files....
- Unit of dispatching
 - An execution paths (may interleaved with other processes)
 - An execution state, dispatching priority.
 - Controlled by OS

Now, What is a thread?

- A thread is an execution path in the code segment
- O.S. provide an individual Program Counter (PC) for each execution path

An example

PC

```
main() {
    // create a TCP/IP socket to use
    s = socket(PF_INET,SOCK_STREAM ,0);

    // bind the server address
    Z = bind(s, (struct sockaddr *)&adr_svrl,
              len_inet);

    // make it a listening socket
    Z = listen(s, 10);

    // start the server loop
    for (;;) {
        // wait for a connect
        c = accept(s, (struct sockaddr *)
                    &adr_clnt,& len_int);
    }
}

PID = fork();
if (PID > 0) {
    // parent process
    close(c);
    continue ;
}
// child process
rx = fdopen(c,"r");
tx = fdopen(dup(c), "w");

// process client's request
.....
fclose(tx); fclose(rx);
exit(0);
```

Comments

- Traditional program is one thread per process.
- The main thread starts with **main()**
- Only one thread (or, program counter) is allowed to execute the code segment
- To add a new PC, you need to **fork()** to have another PC to execute **in another process address space.**

Multithreading

- The ability of an OS to support multiple threads of execution within a single process (**many program counters in a code segment**)
- Windows support multithreading earlier (mid 1990)
- SunOS, Linux came late

The example again

```
main() {  
    // create a TCP/IP socket to use  
    s = socket(PF_INET,SOCK_STREAM ,0);  
  
    // bind the server address  
    Z = bind(s, (struct sockaddr *)&adr_svr,  
             len_inet);  
  
    // make it a listening socket  
    Z = listen(s, 10);  
  
    // start the server loop  
    for (;;) {  
        // wait for a connect  
        c = accept(s, (struct sockaddr*)  
                   &adr_clnt,& len_int);
```

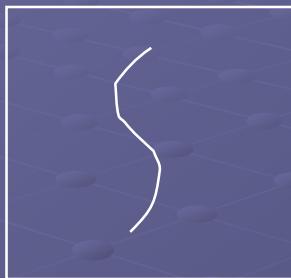


```
hThrd = createThread(Threadfunc,...)  
close(c);  
continue ;  
}  
}  
ThreadFunc() {  
    // child process  
    rx = fdopen(c,"r");  
    tx = fdopen(dup(c), "w");  
  
    // process client's request  
    .....  
    fclose(tx); fclose(rx);  
    exit(0);  
}
```

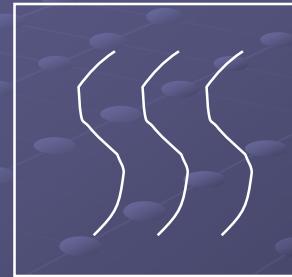
Most APPs are Multithreaded NOW!!

Windows Task Manager Screenshot								
處理程序	效能	應用程式歷程記錄	開機	使用者	詳細資料	服務		
名稱	PID	狀態	使用者名稱	CPU	記憶體 (私人工作集)	執行緒	描述	
POWERPNT.EXE	248	執行中	oolab	00	131,956 K	34	Microsoft PowerPoint	
svchost.exe	1160	執行中	SYSTEM	00	83,576 K	23	Windows Services 的主機處理程序	
HTCVRMarketplac...	7416	執行中	sean9	00	62,804 K	17	HTCVRMarketplaceUserContextHelper	
MsMpEng.exe	744	執行中	SYSTEM	00	61,260 K	37	Antimalware Service Executable	
HTCVRMarketplac...	11364	執行中	oolab	00	45,744 K	17	HTCVRMarketplaceUserContextHelper	
Microsoft.Photos.exe	16016	已暫止	oolab	00	44,464 K	43	Microsoft.Photos.exe	
mspaint.exe	13664	執行中	oolab	00	39,528 K	5	小畫家	
chrome.exe	16520	執行中	oolab	00	29,196 K	43	Google Chrome	
Acrobat.exe	14128	執行中	oolab	00	28,180 K	14	Adobe Acrobat 9.0	
explorer.exe	2676	執行中	oolab	00	25,244 K	67	Windows 檔案總管	
Taskmgr.exe	13796	執行中	oolab	00	21,780 K	21	Task Manager	
svchost.exe	1168	執行中	SYSTEM	00	19,000 K	69	Windows Services 的主機處理程序	
SkypeBrowserHost...	15200	執行中	oolab	00	16,728 K	20	Skype Browser Host	
SearchIndexer.exe	1216	執行中	SYSTEM	00	16,652 K	38	Microsoft Windows Search 索引子	
Skype.exe	11640	執行中	oolab	00	16,508 K	54	Skype	
nw.exe	9952	執行中	sean9	00	15,928 K	16	nw.exe	
svchost.exe	1768	執行中	LOCAL SE...	00	14,964 K	28	Windows Services 的主機處理程序	
chrome.exe	17512	執行中	oolab	00	14,904 K	14	Google Chrome	
dwm.exe	6048	執行中	DWM-2	00	12,052 K	11	桌面視窗管理員	
explorer.exe	3560	執行中	sean9	00	11,908 K	51	Windows 檔案總管	
TortoiseProc.exe	8692	執行中	oolab	00	10,152 K	2	TortoiseSVN client	
chrome.exe	9228	執行中	oolab	00	8,080 K	18	Google Chrome	
RuntimeBroker.exe	6024	執行中	sean9	00	7,664 K	21	Runtime Broker	
Vive.exe	7848	執行中	sean9	00	7,544 K	50	Vive	
svchost.exe	720	執行中	SYSTEM	00	7,448 K	28	Windows Services 的主機處理程序	
ShellExperienceHo...	11228	執行中	oolab	00	7,096 K	25	Windows Shell Experience Host	
TSVNCache.exe	16092	執行中	oolab	00	7,036 K	12	TortoiseSVN status cache	
svchost.exe	1764	執行中	LOCAL SE...	00	6,736 K	17	Windows Services 的主機處理程序	
NisSrv.exe	4344	執行中	LOCAL SE...	00	6,180 K	11	Microsoft Network Realtime Inspection Service	
ApplicationFrameH...	10816	執行中	oolab	00	6,020 K	11	Application Frame Host	
RuntimeBroker.exe	10464	執行中	oolab	00	5,948 K	10	Runtime Broker	

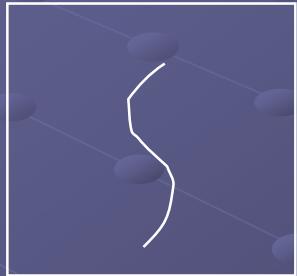
Possible combination of thread and processes



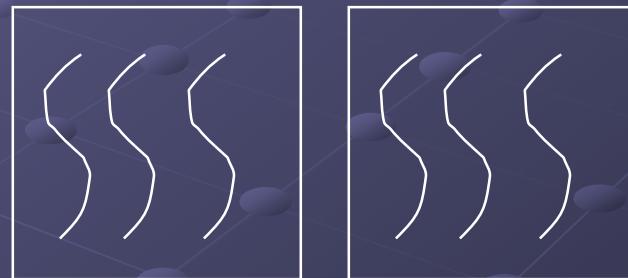
One process one thread



One process multiple thread



Multiple processes
One thread per process



Multiple processes multiple
Threads per process

Process is still there, what's new for thread?

With process

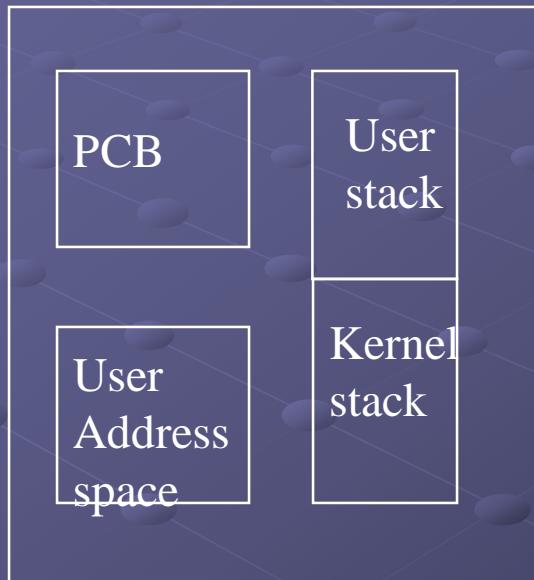
- Virtual address space (holding process image)
- Protected access to CPU, files, and I/O resources

With thread (each thread has its own..)

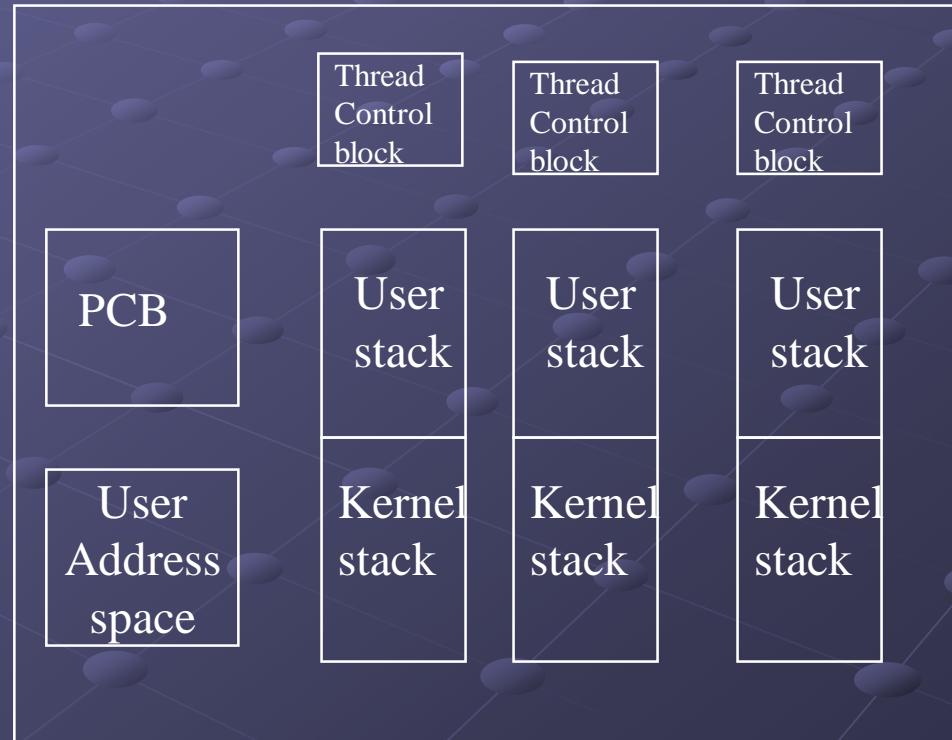
- Thread execution state
- Saved thread context (an independent PC within a process)
- An execution stack
- Per-thread static storage for local variable
- Access to memory and resource of its process, shared with all other threads in that process

Single threaded and multithreaded model

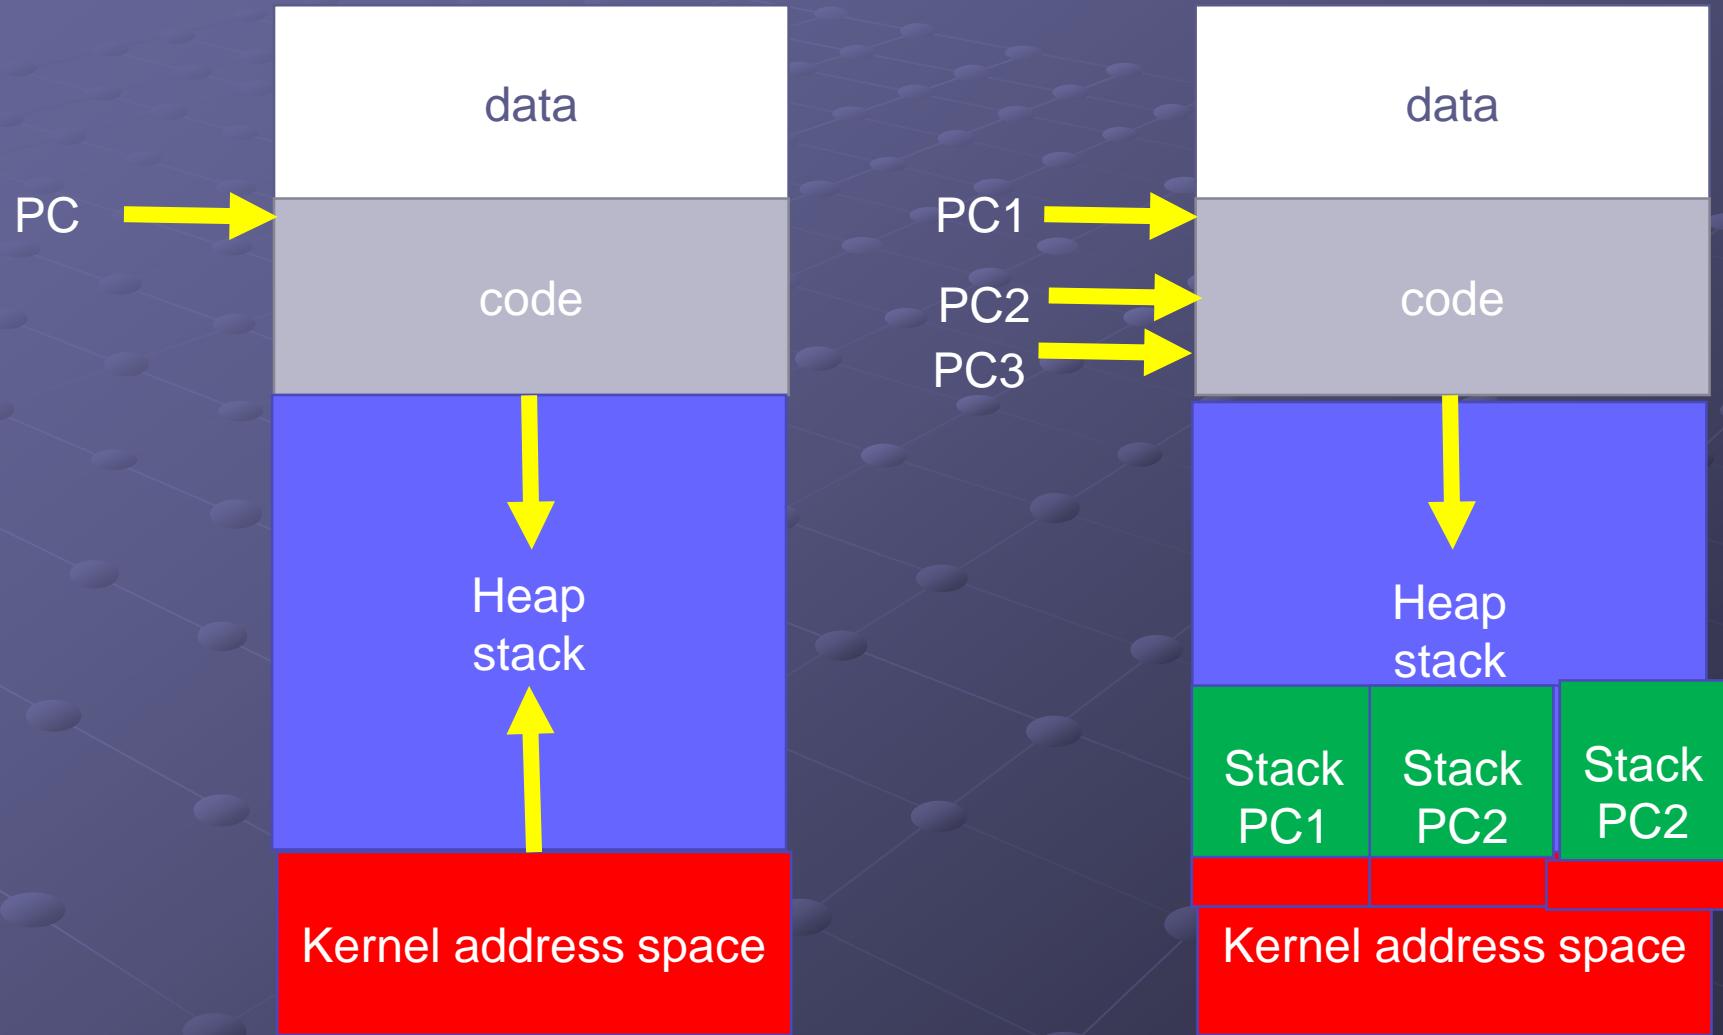
Single-threaded process



Multithread Process model



The meaning of each thread has its own stack but share address space?



C# Multithread examples

All examples assume the following namespaces are imported:

```
using System;
using System.Threading;
```

```
class ThreadTest
{
    static void Main()
    {
        Thread t = new Thread (WriteY);           // Kick off a new thread
        t.Start();                                // running WriteY()

        // Simultaneously, do something on the main thread.
        for (int i = 0; i < 1000; i++) Console.Write ("x");
    }

    static void WriteY()
    {
        for (int i = 0; i < 1000; i++) Console.Write ("y");
    }
}
```

```
xxxxxxxxxxxxxxxxxxxxyyyyyyyyyyyyyyyyyyyyyyyyyyyyyyyyyyyy  

xxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxyyyyyyyyyy  

yyyyyyyyyyyyyyyyyyyyyyyyyyyyyyyyyyyyxxxxxxxxxxxxxxxxxxxx  

xxxxxxxxxxxxxxxxxxxxxxxxxxxxyyyyyyyyyyyyyyyyyyyyyyyyyy  

yyyyyyyyyyyyyyyyxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxx
```

Each thread has its own copy of stack.

```
static void Main()
{
    new Thread (Go).Start();          // Call Go() on a new thread
    Go();                            // Call Go() on the main thread
}

static void Go()
{
    // Declare and use a local variable - 'cycles'
    for (int cycles = 0; cycles < 5; cycles++) Console.Write ('?');
}

???????????
```

Threads that can share on same object

```
class ThreadTest
{
    bool done;

    static void Main()
    {
        ThreadTest tt = new ThreadTest();    // Create a common instance
        new Thread (tt.Go).Start();
        tt.Go();
    }

    // Note that Go is now an instance method
    void Go()
    {
        if (!done) { done = true; Console.WriteLine ("Done"); }
    }
}
```

Because both threads call `Go()` on the same `ThreadTest` instance, they share the `done` field. This results in "Done" being printed once instead of twice:

Done

Static variables are in data segment

Static fields offer another way to share data between threads. Here's the same example with `done` as a static field:

```
class ThreadTest
{
    static bool done;      // Static fields are shared between all threads

    static void Main()
    {
        new Thread (Go).Start();
        Go();
    }

    static void Go()
    {
        if (!done) { done = true; Console.WriteLine ("Done"); }
    }
}
```

Both of these examples illustrate another key concept: that of thread safety (or rather, lack of it!) The output is actually indeterminate: it's possible (though unlikely) that "Done" could be printed twice. If, however, we swap the order of statements in the `Go` method, the odds of "Done" being printed twice go up dramatically:

```
static void Go()
{
    if (!done) { Console.WriteLine ("Done"); done = true; }
}

Done
Done  (usually!)
```

The problem is that one thread can be evaluating the `if` statement right as the other thread is executing the `WriteLine` statement—before it's had a chance to set `done` to true.

Thread lock (mutex) in C# (not IPC, it is faster)

The remedy is to obtain an exclusive lock while reading and writing to the common field. C# provides the lock statement for just this purpose:

```
class ThreadSafe
{
    static bool done;
    static readonly object locker = new object();

    static void Main()
    {
        new Thread (Go).Start();
        Go();
    }

    static void Go()
    {
        lock (locker)
        {
            if (!done) { Console.WriteLine ("Done"); done = true; }
        }
    }
}
```

Key benefits of multithreading

- Less time to create a thread than a process
- Less time to terminate a thread than a process
- Less time to switch a thread
- **Enhance efficiency in communication: no need for kernel to intervene**
- MACH shows a factor of 10

Thread Synchronization vs Process Synchronization

A Comparison of Locking Constructs

Construct	Purpose	Cross-process?	Overhead*
<code>lock (Monitor.Enter / Monitor.Exit)</code>	Ensures just one thread can access a resource (or section of code) at a time	-	20ns
<code>Mutex</code>		Yes	1000ns
<code>SemaphoreSlim</code> (introduced in Framework 4.0)	Ensures not more than a specified number of concurrent threads can access a resource	-	200ns
<code>Semaphore</code>		Yes	1000ns
<code>ReaderWriterLockSlim</code> (introduced in Framework 3.5)	Allows multiple readers to coexist with a single writer	-	40ns
<code>ReaderWriterLock</code> (effectively deprecated)		-	100ns

*Time taken to lock and unlock the construct once on the same thread (assuming no blocking), as measured on an Intel Core i7 860.

Boosted by SMP

- If you install SMP motherboard and SMP enhanced O.S. kernel (Windows and Linux both support)
- Each thread can be assign to an individual CPU.
- The result – boosted performance

Like a global
variable in C,
C++

```
class javathread extends Thread
{
int _threadindex ;
int x;
static int threadno ;
// constructor
javathread(int threadindex) {
    _threadindex = threadindex ;
    threadno ++ ;
}
run() {
    for (int i; i<10000; i++)
        x ++ ;
    System.out.println
    ("this is thread " +
    _threadindex + "of"
    threadno
}
```

Java threads

```
void main() {
    Thread t1 = new javathread(1);
    Thread t2 = new javathread(2);
    t1.start();
    t2.start();
    System.out.println("This is
main thread\n");
}
```

POSSIBLE OUTPUT:

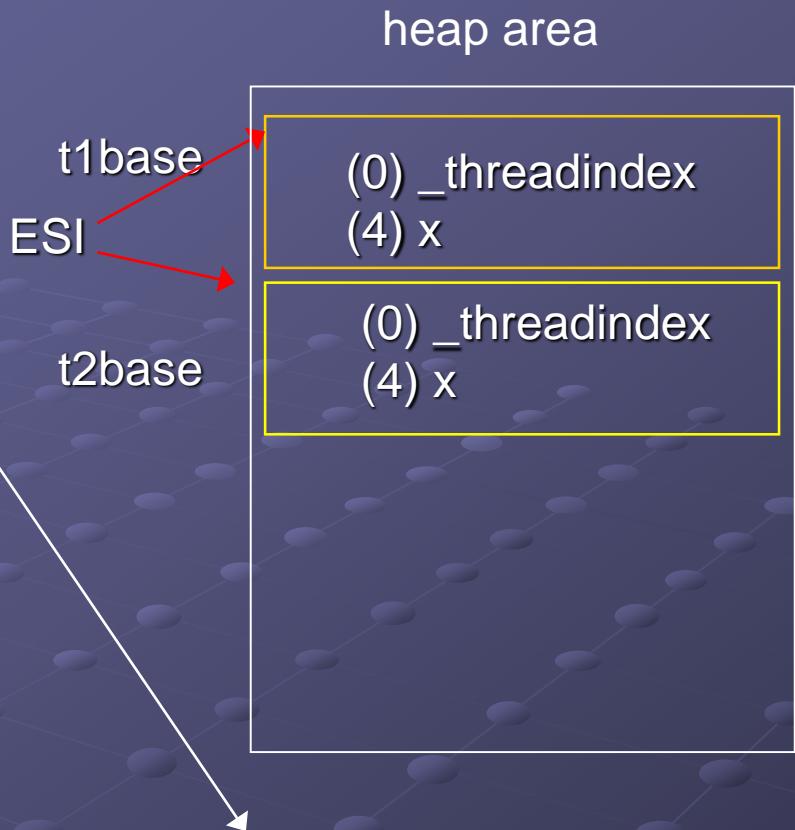
This is main thread !
This is thread 1 of 2
This is thread 2 of 2

code
segment

```
javathread(int threadindex) {  
    _threadindex = threadindex ;  
    threadno ++ ;  
}  
  
run() {  
    for (int i; i<10000; i++)  
        x ++ ;  
    System.out.println  
    ("this is thread " + _threadindex +  
     "of" + threadno );  
}  
  
void main() {  
    Thread t1 = new javathread(1);  
    Thread t2 = new javathread(2);  
    t1.start();  
    t2.start();  
}
```

int threadno ;

data
segment



Win32 Thread – an example (Low level system calls of Windows)

```
int main() {
    HANDLE hThrd ;
    DWORD threadid ;
    int i ;
    for (i=0;i<5;i++) {
        hThrd = CreateThread(NULL,0,ThreadFunc,(LPVOID)i,0,&threadid);
        if (hThr) printf("Thread launched %d\n",i);
    }
    Sleep(2000);
    return EXIT_SUCCESS;
}
DWORD WINAPI ThreadFunc(LPVOID n) {
    int i ;
    for (i=0;i<10;i++) {
        printf("%d%d%d%d%d%d%d%d\n", n,n,n,n,n,n,n,n);
    }
}
```

possible output

0000000

Thread lauched

1111111

1111111

1111111

1111111

2222222

2222222

2222222

2222222

Thread lauched

0000000

Thread lauched

1111111

1111111

1111111

1111111

2222222

2222222

2222222

2222222

Thread lauched

Context
switch may occurs
in the middle
of printf command
- a type of race condition

3333333

3333333

3333333

333344444444

4444444

4444444

4444444

3333

3333333

3333333

3333333

NOTES

- In multi-threading, typically you have no control of output sequences or **scheduling**
- Context switching may occur in any time
 - the probability may be **1/1000000** but consider how many instruction CPU may execute per second
- Your program become concurrent and **nondeterministic**.
 - Same input to a concurrent program may generate different output
 - beware of the probe effect

race condition

Thread 1

.....
AddHead(list,B)
.....

```
AddHead(struct List *list,  
        struct Node *node){  
    node->next = list->head ;  
    list->head = node;  
}
```

Thread 2

.....
AddHead(list,C)
.....

race condition- if you think you are smart enough, things still go wrong

```
AddHead(struct List *list,  
        struct Node *node){  
    while (flag !=0) ;  
    flag = 1 ;  
    node->next = list->head ;  
    list->head = node;  
    flag = 0;  
}
```

```
xor eax, eax ;  
; while (flag !=0)  
L86:  
    cmp _flag, eax  
    jne L86  
; flag = 1  
    mov eax, _list$[esp-4]  
    mov ecx, _node$[esp-4]  
    mov _flag, 1  
; node->next = list->head  
    mov edx, [eax]  
    mov [ecx], edx  
; list->head  
    mov [eax] , ecx  
; flag = 0  
    mov _flag,0
```

What is C runtime library – multi-threaded version

- Original C runtime library is not suitable for multithreaded program
- Original C runtime library uses many global variables and static variables -> cause multithreads to race

NOTES

- conventional C library is not made to execute under multithreaded program
- choose the right C runtime library to link**
 - /ML Single-threaded
 - /MT Multi-threaded (static)
 - /MD Multi-threaded DLL
 - /MLd Debug Single-threaded
 - /MTd Debug Multithreaded (static)
 - /MDd Debug Multithreaded (DLL)

Some suggestions

- avoid using global variables among threads
- do not share GDI objects between thread
- make sure you know the status of the threads you create, do not exit without waiting your threads to terminate
- have the main thread to handle UI (user interface)

Non Multithreading-Safe Code

- When two or more threads enter your code, your code function incorrectly
 - Using global (static) variables (races !!) but programmers are not aware of it
 - Two threads share the same instance of an object (races !!) but programmers are not aware of it.

Synchronization Problems

- There are kinds multithreaded applications which do not need synchronizations among threads.
 - e.g., **telnet** and **bbs** daemon.
 - if your threads do not need any kind of synchronization, such as telnet, you are lucky to get away with **sync horror**.
- Once you need synchronization, you invite concurrency errors to your program.
 - **deadlock**
 - **race conditions**
 - **starvation**

The Path of NO returns

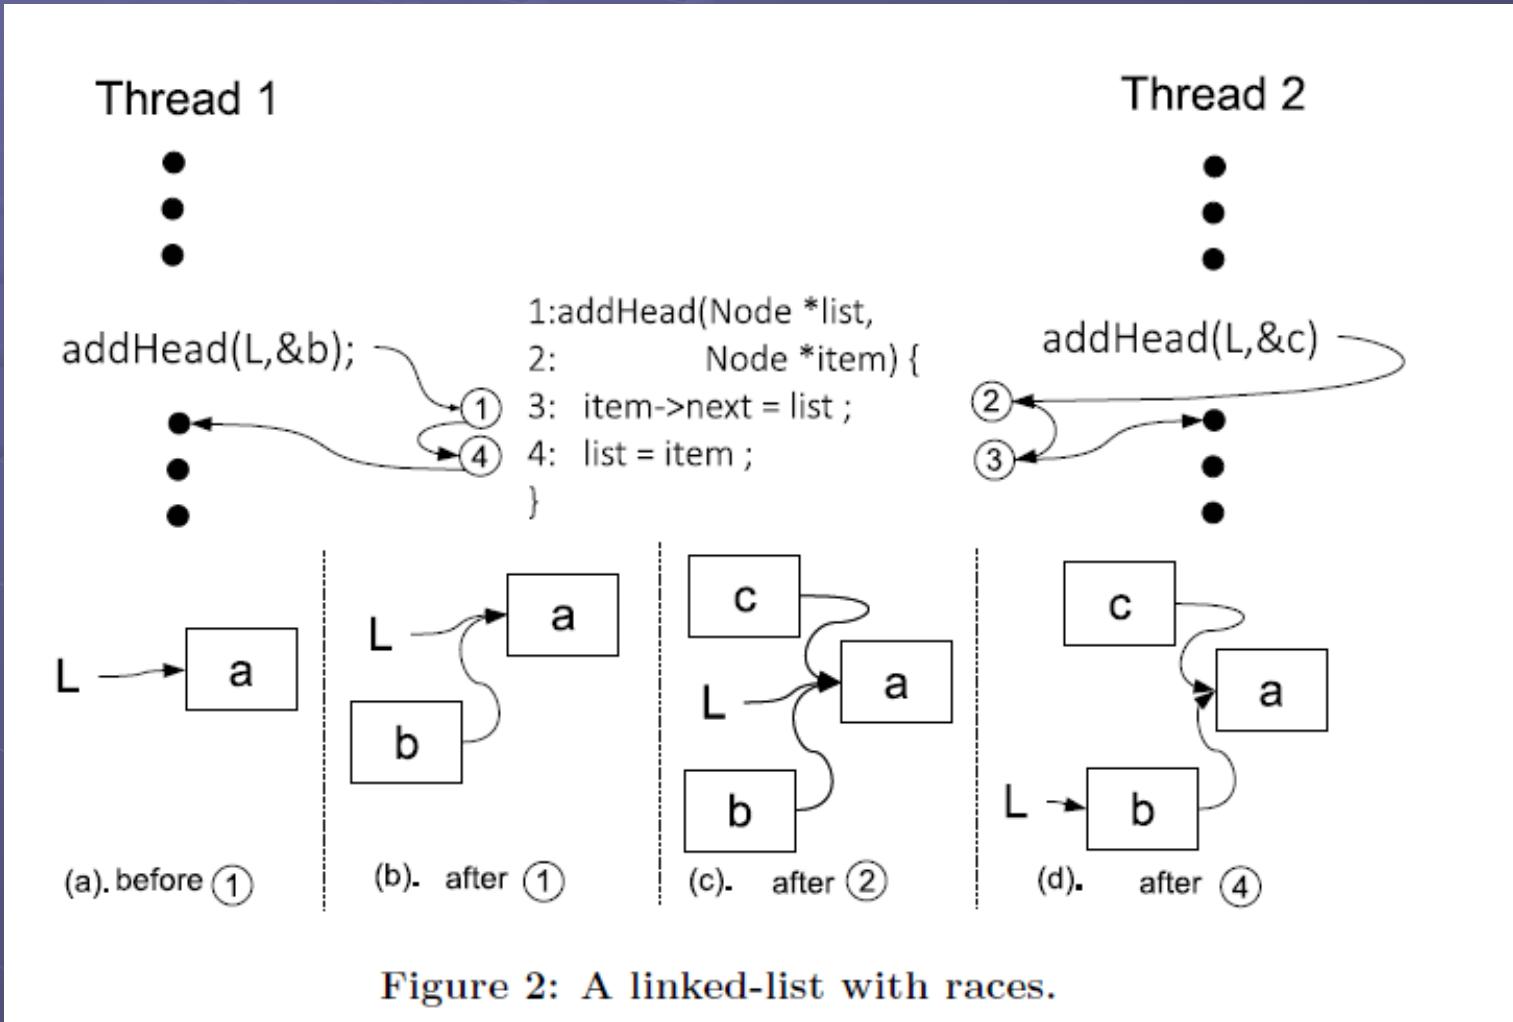
- You have races !!
 - You add synchronizations and thread coordination to fix races !!
 - You introduce deadlock, live lock, starvation...



Concurrency errors

- are known to be
 - hard to detect (irreproducible)
 - hard to debug (probe effect)
 - hard to get a correct fix
- Currently, there is no effective way to detect, debug
- Your better choice is
 - be careful at the beginning (spend more time in design stage)
 - use the simplest synchronization structure that cannot be wrong or proved to be correct but not so efficient
 - do not use complicated synchronization structure unless you are sure what you are doing
 - sacrifices some performance in exchange of correctness.

Why concurrency errors are hard to reproduce, debug, and fix?



Total interleaving permutation

- 1423
- 1234
- 1243
- 2143
- 2134
- 2314
-
- Typically exponential explosion !!
- You may finding a bug-triggering interleaving but cannot reproduce it deterministically because you have no control on O.S. scheduling and context switching
- Yea, we do have a newly developed secrete weapon. In the future, if you encounter concurrency error testing and debugging problem. Please come back to Prof. Cheng.

A Historical Review of Synchronization

Timesharing multiprocesses

Shared memory synchronization problem
Arises
Support hardware sync instruction

The concept of Monitor
Never supported by O.S. system calls



O.S. supports InterProcess Communications

- semaphore (mutex lock)
- message queue
- pipe

Concurrency bugs are nightmare
Concurrency design is in general too difficult for programmers
Only kernel people play with them

Multithreading

- new thread synchronizations
- Monitor are supported in programming language (java)
And usually with a new name
Called synchronized object
- Let the expert do the synchronization In side the shared object
- Clients just use them without knowing the synchronization details

Using semaphores for solving critical section problems

- For n processes
- Initialize S.count to 1
- Then only 1 process is allowed into CS (mutual exclusion)
- To allow k processes into CS, we initialize S.count to k

```
Process Pi:  
repeat  
    P(S);  
    CS  
    V(S);  
    RS  
forever
```

Problems with semaphores

- semaphores provide a powerful tool for enforcing mutual exclusion and coordinate processes
- But P(S) and V(S) are scattered among several processes. Hence, difficult to understand their effects
- Usage must be correct in all the processes
- One bad (or malicious) process can fail the entire collection of processes

IMPORTANT:Suggestion and Trend

- The current trend of synchronization is called **synchronized shared objects**
- gather all the semaphore, mutex, critical sections in an shared object's methods
- Clients need not worry about the synchronization details inside the synchronized shared objects.
- Let the shared object implementer (usually an experienced expert in concurrent programming) worry about the synchronization

Synchronization via a shared object in Java

```
class PRODUCER  
extends Thread {  
buffer b ; // the shared object  
PRODUCER(buffer _b)  
{ b = _b ; }  
run() {  
    do {  
        data = generate_a_data();  
        b.append(data);  
    } while (1);  
}
```

```
class buffer {  
    ?  
}
```

```
class CONSUMER  
extends Thread {  
buffer b ; // the shared object  
CONSUMER(buffer _b)  
{ b = _b ; }  
run() {  
    do {  
        b.take(data);  
        output data ;  
    } while (1);  
}
```

```
main() {  
    buffer rb = new buffer() ;  
    Thread p = new PRODUCER(rb);  
    Thread c = new CONSUMER(rb);  
    p.start();  
    c.start();
```

Let's do it with hand



Appendix – Win32 Synchronization

- The followings are the Win32 synchronization.
- New thread synchronization please refer to the books.

CloseHandle

```
int main() {
    HANDLE hThrd ;
    DWORD threadid ;
    int i ;
    for (i=0;i<5;i++) {
        hThrd =
            CreateThread(NULL,0,ThreadFunc,(LPVOID)i,0,&threadid);
        if (hThr) {
            printf("Thread launched %d\n",i);
            CloseHandle(hThrd)
        }
        Sleep(2000);
    return EXIT_SUCCESS;
}
```

CloseHandle

- Thread is a kernel object
- A kernel object can be owned by several owners
- Each thread when created is owned by two owner- **the creator and the thread itself**. So, the reference count is two.
- **When reference count is down to zero, windows** destroy the object.
- if a process generate many threads but do not close the handle, this process may own many kernel thread object -> **resource leaks**

Misc Win32 function

- **BOOL GetExitCodeThread**

- A non-blocked system call to acquire the status of a thread

- **VOID ExitThread**

- **DWORD WaitForSingleObject**

- avoid busy waiting
 - e.g. WaitForSingleObject(hThrd, INFINITE);

- **DWORD WaitForMultipleObjects**

- same as WaitForSingleObject

Synchronization mechanism in Win32

- Critical sections
- Mutexes
- Semaphore

Critical Sections

- VOID InitializeCriticalSection
- VOID DeleteCriticalSection
- VOID EnterCriticalSection
- VOID LeaveCriticalSection

```
typedef struct _Node {  
    struct _Node *next ;  
    int data ;  
} Node ;  
typedef struct _List {  
    Node *head ;  
    CRITICAL_SECTION  
    critical_sec ;  
} List ;
```

```
List *CreateList() {  
    List *plist = malloc(sizeof(pList));  
    pList-> head = NULL ;  
    InitializeCriticalSection(&pList->critical_sec);  
    return pList ;  
}
```

```
AddHead(struct List *plist,  
        struct Node *node){  
    EnterCriticalSection(&plist->critical_sec);  
    node->next = list->head ;  
    list->head = node;  
    LeaveCriticalSection(&plist->critical_sec);  
}
```

Thread 1

.....
AddHead(list,A)
.....

Thread 2

.....
AddHead(list,B)
.....

```
AddHead(struct List *plist,  
        struct Node *node){  
    EnterCriticalSection(&plist->critical_sec);  
    node->next = list->head ;  
    list->head = node;  
    LeaveCriticalSection(&plist->critical_sec);  
}
```

Continued

- There is at most one thread can manipulate the linked list
- do not use **sleep()** or **Wait..()** in a critical section -> cause deadlock
- if a thread enter a critical section and exit, the critical section is locked by Window NT

Deadlock

```
void SwapLists(List *list1, List *list2) {  
    List *tmp_list ;  
    EnterCriticalSection(list1-> critical_sec);  
    EnterCriticalSection(list2->critical_sec);  
    tmp_list = list1->head  
    list1->head = list2->head  
    list2->head = tmp_list ;  
    LeaveCriticalSection(list1->critical_sec);  
    LeaveCriticalSection(list2->critical_sec);  
}
```

- Holding a resource and then request another resource
-> possible deadlock

Thread 1

Swap_lists

```
void SwapLists(List *list1, List *list2) {  
    List *tmp_list ;  
    EnterCriticalSection(list1-> critical_sec);  
    EnterCriticalSection(list2->critical_sec);  
    tmp_list = list1->head  
    list1->head = list2->head  
    list2->head = tmp_list ;  
    LeaveCriticalSection(list1->critical_sec);  
    LeaveCriticalSection(list2->critical_sec);  
}
```

Thread 2

Swap_lists

Mutex (MUTual EXclusion)

- Similar to critical section
- cause 100 times of time to lock a process/thread
 - CS is executed in user mode (only within a process)
 - mutex is executed in kernel mode , it can cross process (interprocess communication)

Comparision of CS/Mutex

- VOID InitializeCriticalSection
- VOID EnterCriticalSection
- VOID LeaveCriticalSection
- VOID DeleteCriticalSection
- CreateMutex
- OpenMutex
- WaitForSingleObject()
- WaitForMultipleObject()
- MsgWaitForMultipleObject
- ReleaseMutex
- CloseHandle

```
typedef struct _Node {  
    struct _Node *next ;  
    int data ;  
} Node ;  
typedef struct _List {  
    Node *head ;  
    HANDLE hMutex ;  
} List ;
```

```
void SwapLists(List *list1, List *list2) {  
    List *tmp_list ;  
    HANDLE arrhandles[2] ;  
    arrhandles[0] = list1-> hMutex ;  
    arrhandles[1] = list2-> hMutex ;  
    WaitForMultipleObjects(2, arrhandles, TRUE, INFINITE);  
    tmp_list = list1->head  
    list1->head = list2->head  
    list2->head = tmp_list ;  
    ReleaseMutex(arrhandles[0]);  
    ReleaseMutex(arrhandles[1]);  
}
```

```
List *CreateList() {  
    List *plist = malloc(sizeof(pList));  
    pList-> head = NULL ;  
    plist->hMutex = CreateMutex(NULL,false,NULL)  
    return pList ;  
}
```

NOTES

- When a **mutex** is created, no thread/process have a lock on it, **it is ready to trigger**
 - So, a call to Wait..() will immediately return, otherwise a call to Wait...() will be blocked.
- mutex** is exactly the same as a semaphore with semaphore value 1.

Semaphore

- ◆ Synchronization tool (provided by the OS) that do not require busy waiting
- ◆ A semaphore S is an integer variable that, apart from initialization, can only be accessed through 2 atomic and mutually exclusive operations:
 - $P(S)$, $\text{down}(S)$, $\text{wait}(S)$
 - $V(S)$, $\text{up}(S)$, $\text{signal}(S)$
- ◆ To avoid busy waiting: when a process has to wait, it will be put in a blocked queue of processes waiting for the same event

Semaphores in Win32

- HANDLE CreateSemaphore
- void WaitForSingleObject (like P())
- BOOL ReleaseSemaphore (like V())
- They are the same as the semaphore taught in O.S.
- use Wait...() to decrease a semaphore or block a thread when its value is zero.

Semaphore implementation

```
P(S):  
    S.count--;  
    if (S.count<0) {  
        block this process  
        place this process in S.queue  
    }
```

```
V(S):  
    S.count++;  
    if (S.count<=0) {  
        remove a process P from S.queue  
        place this process P on ready list  
    }
```

S.count must be initialized to a nonnegative value
(depending on application)